



# StrokeDetect

---

A Deep Learning Tennis Stroke  
Classification Algorithm

# Introductory Question

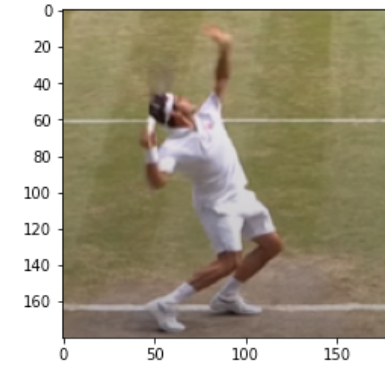
If I have an image of a tennis player in his/her shot motion, how would I detect what stroke they are hitting? Is it possible to automate this detection process?




[Source](#)

# Approach: Image Recognition


- Since tennis strokes are considered image data, we can recognize them via a convolutional neural network
- To keep the problem simple, let's consider three tennis strokes:
  - Forehand (includes volleys)
  - Backhand (includes volleys)
  - Serve
- We can collect images from a uniform dataset of tennis shots and label them as one of the above (3 classes)
- It would then be possible to implement a deep learning network with TensorFlow 2.3, which can be trained to distinguish between forehand, backhand, and serve.



Serve



Forehand



Backhand

# Data Collection



I STARTED BY LOOKING FOR  
PREMADE TENNIS SHOT DETECTION  
DATASETS ON KAGGLE, BUT I  
COULDN'T FIND ANYTHING



I CURATED MY OWN DATASET  
FROM THE 2019 WIMBLEDON  
CHAMPIONSHIPS, SPECIFICALLY THE  
FINAL MATCH BETWEEN FEDERER  
AND DJOKOVIC



THE DATASET I ASSEMBLED  
CONTAINS 334 IMAGES WITH  
ROUGHLY 110 PER CLASS



SMALLER DATASET MAY HAVE LED  
TO A DEGREE OF OVERFITTING  
(WHICH WE WILL SEE LATER)

# Data Preprocessing (Pt. 1)

- After I collected the 334 images, I separated them into different folders based on shot.
- The next step was to set the batch size (number of samples viewed before model updates) and standardize the image height and width:

Set the batch size, image height and width

```
[ ] batch_size = 32
    img_height = 180
    img_width = 180
    data_dir = '/content/gdrive/My Drive/tennis'
```

# Data Preprocessing (Pt. 2)

- The next step is to setup the training and validation datasets. I used a split of 0.8 since it is good practice.
- The batch size was chosen to be 32 since it is good practice for a smaller dataset with less than 100 epochs of training

## Setup validation dataset

```
[ ] val_ds = tf.keras.preprocessing.image_dataset_from_directory(
    data_dir,
    validation_split=0.2,
    subset="validation",
    seed=123,
    image_size=(img_height, img_width),
    batch_size=batch_size)
```

☞ Found 334 files belonging to 3 classes.  
Using 66 files for validation.

## Setup the training dataset

```
[ ] train_ds = tf.keras.preprocessing.image_dataset_from_directory(
    data_dir,
    validation_split=0.2,
    subset="training",
    seed=123,
    image_size=(img_height, img_width),
    batch_size=batch_size)
```

☞ Found 334 files belonging to 3 classes.  
Using 268 files for training.



# Data Preprocessing (Pt. 3)




serve



backhand



forehand





serve



serve



serve



We can now visualize our dataset  
with matplotlib:

# Model Development (Pt. 1)

- After cleaning and visualizing our training data, I developed my network
- I used TensorFlow's Keras API and its sequential model
- We start by defining a normalization layer and setting up our batches of images and labels.
- Pixel values have already been normalized to between 0 and 1

```
[ ] from tensorflow import keras
    from tensorflow.keras import layers
    from tensorflow.keras.models import Sequential

    normalization_layer = layers.experimental.preprocessing.Rescaling(1./255)
```

```
[ ] normalized_ds = train_ds.map(lambda x, y: (normalization_layer(x), y))
    image_batch, labels_batch = next(iter(normalized_ds))
    first_image = image_batch[0]
    # Notice the pixels values are now in `[0,1]`.
    print(np.min(first_image), np.max(first_image))
```

```
0.028096406 1.0
```



# Model Development (Pt. 2)

- We now define the main sequential model.
- Our model consists of three convolution blocks with a max pool layer in each of them.
- There's a fully connected layer with 128 units on top of it and a ReLU activation function.

```
[ ] num_classes = 3

model = Sequential([
    layers.experimental.preprocessing.Rescaling(1./255),
    layers.Conv2D(16, 3, padding='same', activation='relu'),
    layers.MaxPooling2D(),
    layers.Conv2D(32, 3, padding='same', activation='relu'),
    layers.MaxPooling2D(),
    layers.Conv2D(64, 3, padding='same', activation='relu'),
    layers.MaxPooling2D(),
    layers.Dropout(0.2),
    layers.Flatten(),
    layers.Dense(128, activation='relu'),
    layers.Dense(num_classes)
])
```

# Model Training (Pt. 1)

- To train the model, we use `tf.GradientTape()` and record loss values with `SparseCategoricalCrossentropy`.
- The optimizer we are using is Adam and the learning rate is 0.0003.
- The metric we are using is `SparseCategoricalAccuracy()`

```
[ ] # Instantiate an optimizer to train the model.
optimizer = keras.optimizers.Adam(learning_rate=0.0003)
# optimizer = keras.optimizers.SGD(learning_rate=1e-3)

# Instantiate a loss function.
loss_fn = keras.losses.SparseCategoricalCrossentropy(from_logits=True)

# Prepare the metrics.
train_acc_metric = keras.metrics.SparseCategoricalAccuracy()
val_acc_metric = keras.metrics.SparseCategoricalAccuracy()
```

```
[ ] @tf.function
def train_step(x, y):
    with tf.GradientTape() as tape:
        logits = model(x, training=True)
        loss_value = loss_fn(y, logits)
        grads = tape.gradient(loss_value, model.trainable_weights)
        optimizer.apply_gradients(zip(grads, model.trainable_weights))
        train_acc_metric.update_state(y, logits)
    return loss_value
```

# Model Training (Pt. 2)

- We can train our model for 24 epochs
- We iterate over batches of the data size and compare our predicted labels to the true labels
  - We add the current batch loss and update during this process
- At the end of every epoch, we display accuracy, loss, and validation accuracy (at the end)

```
import time

# Keep results for plotting
train_loss_results = []
train_accuracy_results = []

epochs = 20
for epoch in range(epochs):
    print("\nStart of epoch %d" % (epoch,))
    start_time = time.time()

    epoch_loss_avg = tf.keras.metrics.Mean()
    epoch_accuracy = tf.keras.metrics.SparseCategoricalAccuracy()

    # Iterate over the batches of the dataset.
    for step, (x_batch_train, y_batch_train) in enumerate(train_ds):
        loss_value = train_step(x_batch_train, y_batch_train)

        # Track progress
        epoch_loss_avg.update_state(loss_value) # Add current batch loss
        # Compare predicted label to actual label
        # training=True is needed only if there are layers with different
        # behavior during training versus inference (e.g. Dropout).
        epoch_accuracy.update_state(y_batch_train, model(x_batch_train, training=True))

    with train_summary_writer.as_default():
        tf.summary.scalar('loss', epoch_loss_avg.result(), step=epoch)
        tf.summary.scalar('accuracy', epoch_accuracy.result(), step=epoch)

    # Display metrics at the end of each epoch.
    train_acc = train_acc_metric.result()
    print("Training acc over epoch: %.4f" % (float(train_acc),))
    print("Epoch {:03d}: Loss: {:.3f}, Accuracy: {:.3%}".format(epoch,
                                                                epoch_loss_avg.result(),
                                                                epoch_accuracy.result()))

    train_loss_results.append(epoch_loss_avg.result())
    train_accuracy_results.append(epoch_accuracy.result())


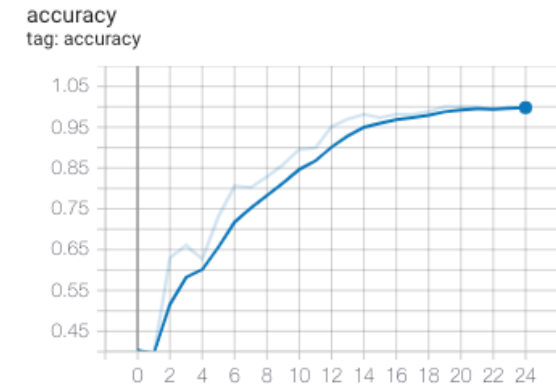
    # Reset training metrics at the end of each epoch
    train_acc_metric.reset_states()
    epoch_loss_avg.reset_states()
    epoch_accuracy.reset_states()

    # Run a validation loop at the end of each epoch.
    for x_batch_val, y_batch_val in val_ds:
        test_step(x_batch_val, y_batch_val)

    val_acc = val_acc_metric.result()
    val_acc_metric.reset_states()
    print("Validation acc: %.4f" % (float(val_acc),))
    print("Time taken: %.2fs" % (time.time() - start_time))
```


# Results (Pt. 1)

- As you can see in the graphs below, our model behavior was correct in that loss decreases accuracy increases over the epochs
- By the final epoch (24), the loss was 0.021 and the accuracy was 100.00%



# Results (Pt. 2)

Training Metrics



# Inference Results

```
test_path = "/content/gdrive/My Drive/test/test2.png"


img1 = PIL.Image.open(test_path)
img1_resized = img1.resize((180, 180))
plt.imshow(img1_resized)

img = keras.preprocessing.image.load_img(
    test_path, target_size=(img_height, img_width)
)
img_array = keras.preprocessing.image.img_to_array(img)
img_array = tf.expand_dims(img_array, 0) # Create a batch

predictions = model.predict(img_array)
score = tf.nn.softmax(predictions[0])

print(
    "This image most likely belongs to {} with a {:.2f} percent confidence."
    .format(class_names[np.argmax(score)], 100 * np.max(score))
)
```

This image most likely belongs to forehand with a 99.04 percent confidence.





# Conclusions and Future Work

- In the future, I hope to add to my project by expanding the dataset.
- When you try inference for shots not from Wimbledon, the performance drops, so I want to add in 250 images at minimum for all the major surfaces (grass, hard, clay)
- I also want to add in more players. Right now, the model is thrown off by left-handed players, so incorporating lefties like Rafael Nadal will be interesting and will require a large dataset.
- To increase the sophistication of the algorithm, I would like to try other existing pre-trained CNNs as well (ResNet, VGG, etc.)